
Understanding Memory Management

Dipesh Kafle

Prerequisites

Memory Layout




Figure 1: A rough classification of program's memory segments

Memory Layout




Figure 1: A rough classification of program's memory segments

- In practice, the stack grows towards lower addresses, the heap towards higher(the diagram has it the other way around, but that doesn't matter).

Memory Layout




Figure 1: A rough classification of program's memory segments

- In practice, the stack grows towards lower addresses, the heap towards higher(the diagram has it the other way around, but that doesn't matter).
- What are all these things ??

- **Code:** Generated target code has a fixed size, allowing it to be stored in a statically determined area called Code, usually at the low end of memory.

Code and Static segments

- **Code:** Generated target code has a fixed size, allowing it to be stored in a statically determined area called Code, usually at the low end of memory.
- **Static:** Statically determined data objects, such as global constants and data generated by the compiler at compile time, can be stored in another area called Static.

Code and Static segments

- **Code:** Generated target code has a fixed size, allowing it to be stored in a statically determined area called Code, usually at the low end of memory.
- **Static:** Statically determined data objects, such as global constants and data generated by the compiler at compile time, can be stored in another area called Static.

```
1  const char* s = "Lorem Ipsum something something";
2  int main(){
3      const char* string_arr[] = {"Made", "with", "love",
   ↪   "by", "Delta", "Force"};
4      return 0;
5  }
```

All the strings used in the above code segment are stored in static section, while the instructions generated for the program will be in code section.

- The stack will store things such as local variables, return address from a function call, etc.

Stack and Stack Allocation(continued)

```
1  int main(){
2  // All of this is
3  // doing stack
   ⇨ allocation
4  int a = 10;
5  int b = 20;
6  int arr[2] = {1,2};
7  return 0;
8  }
```

Stack and Stack Allocation(continued)

```
1  int main(){  
2  // All of this is  
3  // doing stack  
   ↪ allocation  
4  int a = 10;  
5  int b = 20;  
6  int arr[2] = {1,2};  
7  return 0;  
8  }
```



Figure 2: Stack Layout for above code

Heap and Heap Allocation

- Many programming languages allow the programmer to allocate and deallocate data under program control.

Heap and Heap Allocation

- Many programming languages allow the programmer to allocate and deallocate data under program control.
- The heap is used to manage long-lived data.

Heap and Heap Allocation

- Many programming languages allow the programmer to allocate and deallocate data under program control.
- The heap is used to manage long-lived data.
- Unavoidable when we want to allocate memory whose size is known only when the program is running(dynamic allocation).

Heap and Heap Allocation

- Many programming languages allow the programmer to allocate and deallocate data under program control.
- The heap is used to manage long-lived data.
- Unavoidable when we want to allocate memory whose size is known only when the program is running(dynamic allocation).
- C/C++ has **malloc/realloc/free** functions for doing heap memory management.

Heap and Heap Allocation

- Many programming languages allow the programmer to allocate and deallocate data under program control.
- The heap is used to manage long-lived data.
- Unavoidable when we want to allocate memory whose size is known only when the program is running(dynamic allocation).
- C/C++ has `malloc/realloc/free` functions for doing heap memory management.

```
1  int* f(int n){
2      return malloc(n*sizeof(int));
3  }
4  int main(){
5      int n ;
6      scanf("%d", &n);
7      int *arr = f(n); // arr is heap allocated, returned from call to f
8      free(arr); //Since, we're good programmers, we'll free the memory as
        ↪ well.
9  }
```

What exactly are malloc/realloc/free?

What exactly are malloc/realloc/free?

- `malloc(x)` : allocate x bytes in heap

What exactly are malloc/realloc/free?

- `malloc(x)` : allocate x bytes in heap
- `realloc(p, x)` : resize previously allocated heap memory

What exactly are malloc/realloc/free?

- `malloc(x)` : allocate x bytes in heap
- `realloc(p, x)` : resize previously allocated heap memory
- `free(p)` : return heap memory to the operating system

Introduction to Memory Management

- Memory management is all about using `heap memory` correctly.

- Memory management is all about using **heap memory** correctly.
- Stack allocations don't need to be freed; they're automatically managed with **scopes** (we'll talk about scopes in the next slide).

Introduction to Memory Management

- Memory management is all about using **heap memory** correctly.
- Stack allocations don't need to be freed; they're automatically managed with **scopes** (we'll talk about scopes in the next slide).
- If it's done incorrectly, the program can **crash** or **slow down**.

Introduction to Memory Management

- Memory management is all about using **heap memory** correctly.
- Stack allocations don't need to be freed; they're automatically managed with **scopes** (we'll talk about scopes in the next slide).
- If it's done incorrectly, the program can **crash** or **slow down**.
- There are different techniques for managing heap memory.

What is this scope thing??

What is this scope thing??

Easier to understand with an example

What is this scope thing??

Easier to understand with an example

```
1  // NOTE: This function won't compile
2  int f(){ // scope '1 starts
3      int a = 10;
4      { // scope '2 starts
5          int b = 20;
6          { // scope '3 starts
7              int c = 50;
8          } // scope '3 ends
9      } // scope '2 ends
10
11     if(a == 10){ // scope '4 starts
12         int c = 30;
13     } // scope '4 ends
14     else { // scope '5 starts
15         int c = 90;
16     } // scope '5 ends
17     return b; // This fails because it's not in scope
18 } // scope '1 ends
```


What is this scope thing?? (continued)




Figure 3: Pictorial representation of scope heirarchy

Why should I care about freeing memory? Is it really a problem?

Why should I care about freeing memory? Is it really a problem?

- If your system has infinite memory, you don't need to worry. However, since memory is finite, you must take care.

Why should I care about freeing memory? Is it really a problem?

- If your system has infinite memory, you don't need to worry. However, since memory is finite, you must take care.
- If one program uses up all the memory, other programs that require memory won't be able to function properly.

Why should I care about freeing memory? Is it really a problem?

- If your system has infinite memory, you don't need to worry. However, since memory is finite, you must take care.
- If one program uses up all the memory, other programs that require memory won't be able to function properly.
- Your program may **crash** if it requests more memory than the operating system can provide.

Why should I care about freeing memory? Is it really a problem?

- If your system has infinite memory, you don't need to worry. However, since memory is finite, you must take care.
- If one program uses up all the memory, other programs that require memory won't be able to function properly.
- Your program may **crash** if it requests more memory than the operating system can provide.
- **Memory leaks** can have a significant impact on long-running programs such as **web servers, editors, and IDEs**.

Why should I care about freeing memory? Is it really a problem?

- If your system has infinite memory, you don't need to worry. However, since memory is finite, you must take care.
- If one program uses up all the memory, other programs that require memory won't be able to function properly.
- Your program may **crash** if it requests more memory than the operating system can provide.
- **Memory leaks** can have a significant impact on long-running programs such as **web servers, editors, and IDEs**.

Important terminology

- **Memory Leak**: It happens when you ask the operating system for memory but don't return it back.

Ways to manage memory

We have two ways to do memory management.

We have two ways to do memory management.

- **Manual Memory Management** : Languages such as C, C++, Rust, etc have this

We have two ways to do memory management.

- **Manual Memory Management** : Languages such as C, C++, Rust, etc have this
- **Automatic Memory Management**: Languages such as Python, Java, Go, JavaScript, Swift, etc have this.

Manual Memory Management

Scenarios where you can go wrong

```
1  // NOTE: this is a dumb example to
   ↪ show where things can go wrong,
2  // I don't actually write code like this
3  int* allocate_and_throw_exn(int n){
4      int *arr = malloc(n*sizeof(int));
5      if (n < 10){
6          throw runtime_error("n < 10");
7      }
8      return arr;
9  }
10 int main(){
11     try {
12         auto *arr =
           ↪ allocate_and_throw_exn(2);
13         free(arr);
14     } catch(const std::runtime_error
           ↪ &e){
15         cout << "Error:" <<e.what() <<
           ↪ endl;
16     }
17 }
```

Scenarios where you can go wrong

```
1  // NOTE: this is a dumb example to
   ↪ show where things can go wrong,
2  // I don't actually write code like this
3  int* allocate_and_throw_exn(int n){
4      int *arr = malloc(n*sizeof(int));
5      if (n < 10){
6          throw runtime_error("n < 10");
7      }
8      return arr;
9  }
10 int main(){
11     try {
12         auto *arr =
13             ↪ allocate_and_throw_exn(2);
14         free(arr);
15     } catch(const std::runtime_error
16             ↪ &e){
17         cout << "Error:" <<e.what() <<
18             ↪ endl;
19     }
20 }
```


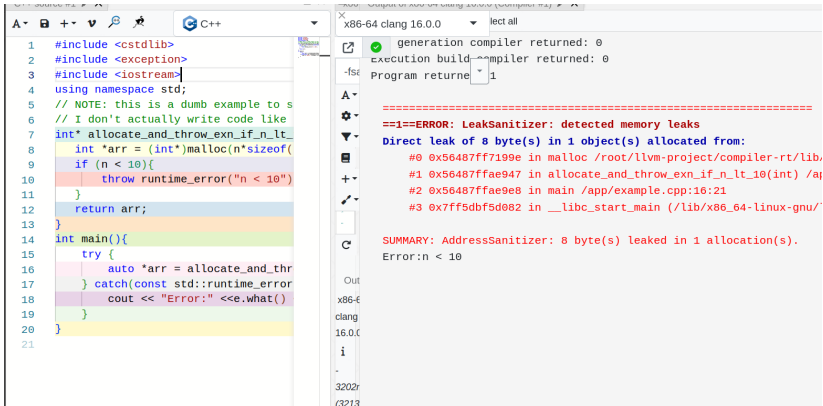


Figure 4: Flow for the leaking code

Scenarios where you can go wrong



The image shows a C++ code editor on the left and its output window on the right. The code in the editor is a simple program that allocates memory and throws an exception if the size is less than 10. The output window shows the program's execution, including the compiler's return code, the program's return code, and the AddressSanitizer's report of a memory leak.

```
1 #include <cstdlib>
2 #include <exception>
3 #include <iostream>
4 using namespace std;
5 // NOTE: this is a dumb example to s
6 // I don't actually write code like
7 int* allocate_and_throw_exn_if_n_lt
8     int *arr = (int*)malloc(n*sizeof(
9     if (n < 10){
10         throw runtime_error("n < 10")
11     }
12     return arr;
13 }
14 int main(){
15     try {
16         auto *arr = allocate_and_thr
17         catch(const std::runtime_error
18             cout << "Error:" <<e.what()
19         }
20     }
21 }
```

Output window (x86_64 clang 16.0.0):

```
generation compiler returned: 0
execution build-compiler returned: 0
Program returned: 1

=====
==1==ERROR: LeakSanitizer: detected memory leaks
Direct leak of 8 byte(s) in 1 object(s) allocated from:
#0 0x56487ff7199e in malloc /root/llvm-project/compiler-rt/lib,
#1 0x56487ffae947 in allocate_and_throw_exn_if_n_lt_10(int) /app
#2 0x56487ffae9e8 in main /app/example.cpp:16:21
#3 0x7ff5dbf5d082 in __libc_start_main (/lib/x86_64-linux-gnu/

SUMMARY: AddressSanitizer: 8 byte(s) leaked in 1 allocation(s).
Error:n < 10
```

Figure 5: Memory Leak Detected by address sanitizer

How to fix it??

- DON'T WRITE DUMB CODE LIKE I DID

- DON'T WRITE DUMB CODE LIKE I DID
- More high level language(than C) like C++, Rust provide us with smart ways to manage memory

- DON'T WRITE DUMB CODE LIKE I DID
- More high level language(than C) like C++, Rust provide us with smart ways to manage memory
- They come built-in with smart pointer types like `unique_ptr` (C++)

Fixing the code with smart pointer

```
1  # LEAKING CODE
2  int* allocate_and_throw_exn(int n){
3      int *arr = malloc(n*sizeof(int));
4      if (n < 10){
5          throw runtime_error("n < 10");
6      }
7      return arr;
8  }
9  int main(){
10     try {
11         auto *arr =
12             ↪ allocate_and_throw_exn(2);
13         free(arr);
14     } catch(const std::runtime_error
15             ↪ &e){
16         cout << "Error:" <<e.what() <<
17             ↪ endl;
18     }
19 }
```

Fixing the code with smart pointer

```
1  # LEAKING CODE
2  int* allocate_and_throw_exn(int n){
3      int *arr = malloc(n*sizeof(int));
4      if (n < 10){
5          throw runtime_error("n < 10");
6      }
7      return arr;
8  }
9  int main(){
10     try {
11         auto *arr =
12             ↪ allocate_and_throw_exn(2);
13         free(arr);
14     } catch(const std::runtime_error
15             ↪ &e){
16         cout << "Error:" <<e.what() <<
17             ↪ endl;
18     }
19 }
```

```
1  # FIXED CODE
2  std::unique_ptr<int[]>
3      ↪ allocate_and_throw_exn(int n){
4      auto arr =
5          ↪ std::unique_ptr<int[]>(new
6          ↪ int[n]);
7      if (n < 10){
8          throw runtime_error("n < 10");
9      }
10     return arr;
11 }
12 int main(){
13     try {
14         auto arr =
15             ↪ allocate_and_throw_exn(2);
16     } catch(const std::runtime_error
17             ↪ &e){
18         cout << "Error:" <<e.what() <<
19             ↪ endl;
20     }
21 }
```

Fixing the code with smart pointer

```
1  # LEAKING CODE
2  int* allocate_and_throw_exn(int n){
3      int *arr = malloc(n*sizeof(int));
4      if (n < 10){
5          throw runtime_error("n < 10");
6      }
7      return arr;
8  }
9  int main(){
10     try {
11         auto *arr =
12             ↪ allocate_and_throw_exn(2);
13         free(arr);
14     } catch(const std::runtime_error
15             ↪ &e){
16         cout << "Error:" <<e.what() <<
17             ↪ endl;
18     }
19 }
```

```
1  # FIXED CODE
2  std::unique_ptr<int[]>
3  ↪ allocate_and_throw_exn(int n){
4      auto arr =
5          ↪ std::unique_ptr<int[]>(new
6          ↪ int[n]);
7      if (n < 10){
8          throw runtime_error("n < 10");
9      }
10     return arr;
11 }
12 int main(){
13     try {
14         auto arr =
15             ↪ allocate_and_throw_exn(2);
16     } catch(const std::runtime_error
17             ↪ &e){
18         cout << "Error:" <<e.what() <<
19             ↪ endl;
20     }
21 }
```

What the hell just happened??

How do smart pointers work?

How do smart pointers work?

- Uses **scope** to track lifetime of a pointer(scopes mentioned in [previous section](#))

How do smart pointers work?

- Uses **scope** to track lifetime of a pointer(scopes mentioned in [previous section](#))
- C++ uses **destructors** to run code when an object goes out of scope.
(due to **RAII** in C++)

What is RAI in C++?

RAII can be summarized as follows:

- encapsulate each resource into a class, where
 - the constructor acquires the resource and establishes all class invariants or throws an exception if that cannot be done,
 - the destructor releases the resource and never throws exceptions;

Figure 6: RAII

Destructors In Action

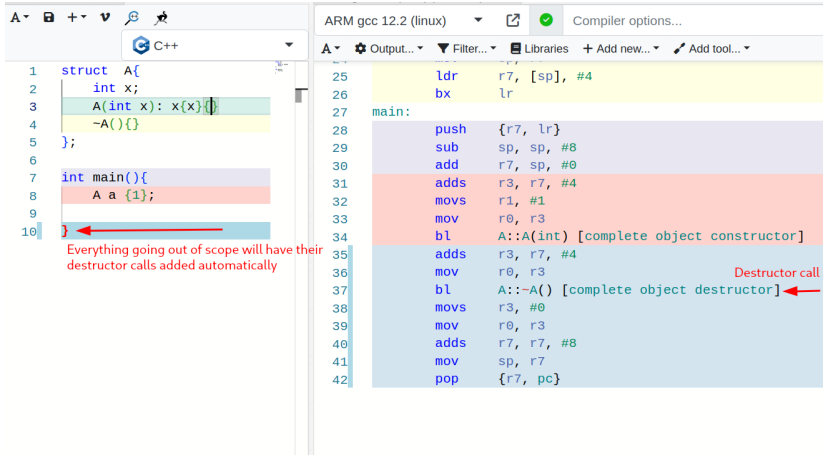


Figure 7: Destructor call added automatically

Let's make our own unique_ptr

Let's make our own unique_ptr

```
1  #include <iostream>
2  using namespace std;
3  class int_ptr{
4      int* x ;
5  public:
6      int_ptr(int x): x{new int(x)}{ }
7      ~int_ptr(){
8          delete x;
9      }
10     int& operator*(){
11         return *x;
12     }
13 };
14 int main(){
15     int_ptr one(1);
16     cout << *one << endl;
17 }
```

```
ASM generation compiler returned: 0
Execution build compiler returned: 0
Program returned: 0
1
```

Figure 8: int_ptr working without any leaks

Is `unique_ptr` the only smart pointer??

Is `unique_ptr` the only smart pointer??

NO

Isn't `unique_ptr` perfect already? Why would I need anything else?

Isn't `unique_ptr` perfect already? Why would I need anything else?

- Problem with `unique_ptr` is that it can have only one owner.

Isn't `unique_ptr` perfect already? Why would I need anything else?

- Problem with `unique_ptr` is that it can have only one owner.

Case in Point

- A `unix` file descriptor

Isn't `unique_ptr` perfect already? Why would I need anything else?

- Problem with `unique_ptr` is that it can have only one owner.

Case in Point

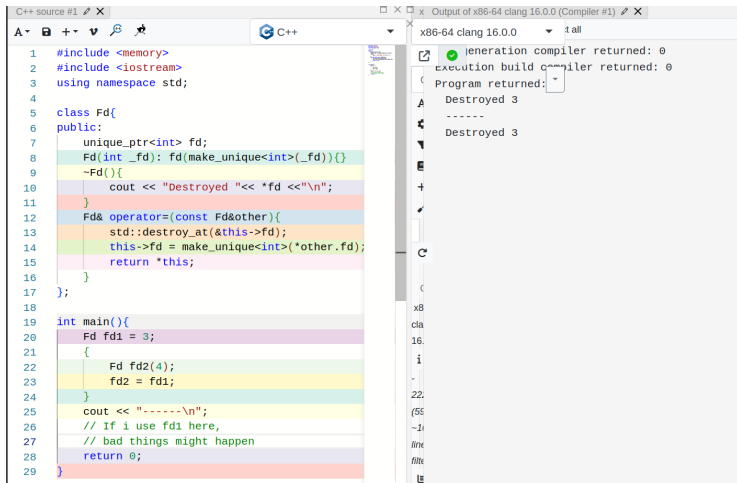
- A `unix file descriptor`
- A file descriptor can have multiple owners. It should only be freed when all the owners go out of scope.

DISCLAIMER: The code that I'm about to show is cursed and no one should model a file descriptor like this

DISCLAIMER: The code that I'm about to show is cursed and no one should model a file descriptor like this

But, we'll do it still(for the purpose of this talk)

File Descriptors (continued)



The screenshot shows a C++ IDE with two panes. The left pane displays the source code for a program that uses `std::unique_ptr` to manage file descriptors. The right pane shows the output of the compiler, indicating successful compilation and execution.

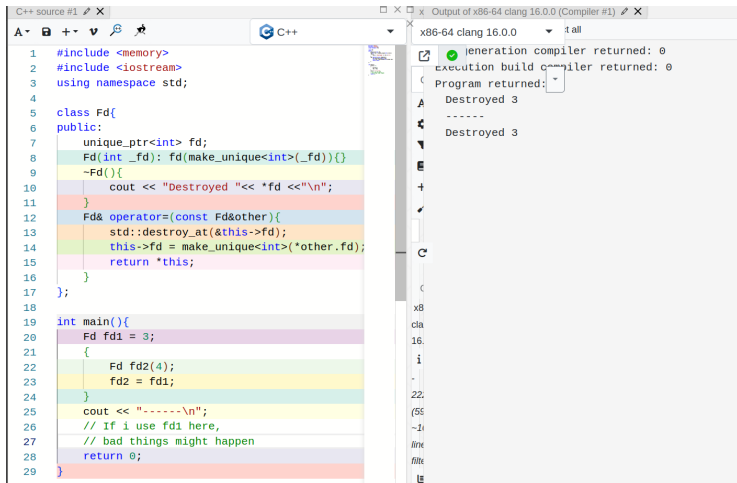
```
1 #include <memory>
2 #include <iostream>
3 using namespace std;
4
5 class Fd{
6 public:
7     unique_ptr<int> fd;
8     Fd(int _fd): fd(make_unique<int>(_fd)){}
9     ~Fd(){
10         cout << "Destroyed "<< *fd << "\n";
11     }
12     Fd& operator=(const Fd&other){
13         std::destroy_at(&this->fd);
14         this->fd = make_unique<int>(*other.fd);
15         return *this;
16     }
17 };
18
19 int main(){
20     Fd fd1 = 3;
21     {
22         Fd fd2(4);
23         fd2 = fd1;
24     }
25     cout << "-----\n";
26     // If i use fd1 here,
27     // bad things might happen
28     return 0;
29 }
```

Output of x86-64 clang 16.0.0 (Compiler #1):

```
x86-64 clang 16.0.0
generation compiler returned: 0
execution build compiler returned: 0
Program returned:
Destroyed 3
-----
Destroyed 3
```

Figure 9: File Descriptor with `unique_ptr`

File Descriptors (continued)



The screenshot shows a C++ IDE with two windows. The left window, titled 'C++ source #1', contains the following code:

```
1 #include <memory>
2 #include <iostream>
3 using namespace std;
4
5 class Fd{
6 public:
7     unique_ptr<int> fd;
8     Fd(int _fd): fd(make_unique<int>(_fd)){}
9     ~Fd(){
10         cout << "Destroyed "<< *fd << "\n";
11     }
12     Fd& operator=(const Fd&other){
13         std::destroy_at(&this->fd);
14         this->fd = make_unique<int>(*other.fd);
15         return *this;
16     }
17 };
18
19 int main(){
20     Fd fd1 = 3;
21     {
22         Fd fd2(4);
23         fd2 = fd1;
24     }
25     cout << "-----\n";
26     // If i use fd1 here,
27     // bad things might happen
28     return 0;
29 }
```

The right window, titled 'Output of x86-64 clang 16.0.0 (Compiler #1)', shows the compiler output:

```
x86-64 clang 16.0.0
generation compiler returned: 0
execution build compiler returned: 0
Program returned:
Destroyed 3
-----
Destroyed 3
```

Figure 9: File Descriptor with unique_ptr

Not possible to model this correctly

Another classic example (why only `unique_ptr` isn't enough?)

- SPMC(Single Producer Multiple Consumer) and MPMC(Multiple Producer Multiple Consumer) problem

Another classic example (why only unique_ptr isn't enough?)

- SPMC(Single Producer Multiple Consumer) and MPMC(Multiple Producer Multiple Consumer) problem



Figure 10: MPMC queue

Another classic example (why only unique_ptr isn't enough?)

- SPMC(Single Producer Multiple Consumer) and MPMC(Multiple Producer Multiple Consumer) problem



Figure 10: MPMC queue

- Requires multiple owners for producer end as well as the consumer end

Introduction to Automatic Memory Management

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

IMPORTANT TERMINOLOGY

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

IMPORTANT TERMINOLOGY

Data that cannot be referenced is generally known as **garbage**

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

IMPORTANT TERMINOLOGY

Data that cannot be referenced is generally known as **garbage**

Two techniques used primarily

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

IMPORTANT TERMINOLOGY

Data that cannot be referenced is generally known as **garbage**

Two techniques used primarily

- We catch the transitions as **reachable**(in-scope) objects turn **unreachable**(out-of-scope).

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

IMPORTANT TERMINOLOGY

Data that cannot be referenced is generally known as **garbage**

Two techniques used primarily

- We catch the transitions as **reachable**(in-scope) objects turn **unreachable**(out-of-scope). Example: **Reference Counting** based memory management in Python, Swift, etc., C++ `shared_ptr`, Rust's `Rc` and `Arc` types

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

IMPORTANT TERMINOLOGY

Data that cannot be referenced is generally known as **garbage**

Two techniques used primarily

- We catch the transitions as **reachable**(in-scope) objects turn **unreachable**(out-of-scope). Example: **Reference Counting** based memory management in Python, Swift, etc., C++ `shared_ptr`, Rust's `Rc` and `Arc` types
- We periodically locate all the **reachable** objects and then infer that all the other objects are **unreachable**.

Introduction to Automatic Memory Management

- The programming language you're using is responsible for memory management. The language runtime handles everything.
- Programmer doesn't have the mental burden of dealing with memory.
- All the allocation/deallocation calls are hidden from the programmer.

IMPORTANT TERMINOLOGY

Data that cannot be referenced is generally known as **garbage**

Two techniques used primarily

- We catch the transitions as **reachable**(in-scope) objects turn **unreachable**(out-of-scope). Example: **Reference Counting** based memory management in Python, Swift, etc., C++ `shared_ptr`, Rust's `Rc` and `Arc` types
- We periodically locate all the **reachable** objects and then infer that all the other objects are **unreachable**. Example: **Trace Based Garbage Collection** such as mark and sweep garbage collector used in language like Java, JavaScript, etc.

Reference Counting

- All the transitions from reachable to unreachable are caught immediately.

- All the transitions from reachable to unreachable are caught immediately.

Backed by simple rules

- All the transitions from reachable to unreachable are caught immediately.

Backed by simple rules

- Every allocated pointer has a **reference count** associated with it

- All the transitions from reachable to unreachable are caught immediately.

Backed by simple rules

- Every allocated pointer has a **reference count** associated with it
- Every copy leads to reference count increment.

- All the transitions from reachable to unreachable are caught immediately.

Backed by simple rules

- Every allocated pointer has a **reference count** associated with it
- Every copy leads to reference count increment.
- Every time a reference of the pointer goes out of scope, reference count is decremented by 1.

- All the transitions from reachable to unreachable are caught immediately.

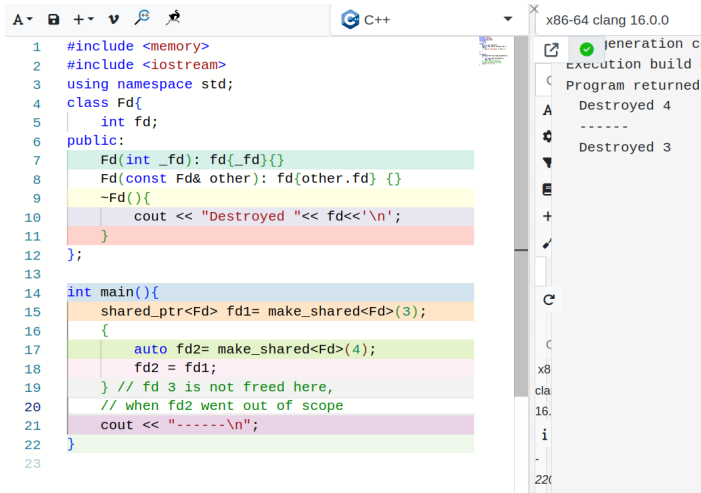
Backed by simple rules

- Every allocated pointer has a **reference count** associated with it
- Every copy leads to reference count increment.
- Every time a reference of the pointer goes out of scope, reference count is decremented by 1.
- When the reference count goes to 0, it is safe to **free** the object.

- It is a **reference-counted pointer**, manages the **reference count** elegantly with assignment operator overload, copy constructor overload, destructors, etc.

Let's solve our file descriptor issue

Let's solve our file descriptor issue



The screenshot shows a C++ IDE with a code editor on the left and a console output on the right. The code defines a class `Fd` with a file descriptor `fd` and implements copy and move constructors and a destructor. The `main` function uses `shared_ptr` to manage `Fd` objects. It creates `fd1` with 3 and `fd2` with 4, then assigns `fd2` to `fd1`, effectively destroying the object with fd 4. The console output shows the program returning, destroyed 4, and destroyed 3.

```
1  #include <memory>
2  #include <iostream>
3  using namespace std;
4  class Fd{
5  |   int fd;
6  public:
7  |   Fd(int _fd): fd{_fd}{}
8  |   Fd(const Fd& other): fd{other.fd} {}
9  |   ~Fd(){
10 |       cout << "Destroyed " << fd << '\n';
11 |   }
12 };
13
14 int main(){
15 |   shared_ptr<Fd> fd1= make_shared<Fd>(3);
16 |   {
17 |       auto fd2= make_shared<Fd>(4);
18 |       fd2 = fd1;
19 |   } // fd 3 is not freed here,
20 |   // when fd2 went out of scope
21 |   cout << "-----\n";
22 }
23
```

Console Output:

```
x86-64 clang 16.0.0
[✓] generation c
Execution build
Program returned
Destroyed 4
-----
Destroyed 3
```

Figure 11: File Descriptor with `shared_pointer`

Trace Based Garbage Collection

- Periodically locates all the **reachable** objects and then infers that all the other objects are **unreachable**

- Periodically locates all the **reachable** objects and then infers that all the other objects are **unreachable**
- A pointer is reachable \Rightarrow There's a root somewhere(in stack, register, static, etc.) from where we can transitively access the pointer.

What's a root?

What's a root?

Think of them as root nodes of a graph, something that you have direct access to.

What's a root?

Think of them as root nodes of a graph, something that you have direct access to.

A **GC** will treat all the stack variables, registers, static section, as roots.

What's a root?

Think of them as root nodes of a graph, something that you have direct access to.

A **GC** will treat all the stack variables, registers, static section, as roots.

```
1  type baz = {  
2      x: int;  
3      y: int  
4  }  
5  type bar = {  
6      a: int;  
7      b: baz  
8  }  
9  type foo = {  
10     field1 : bar;  
11     field2: int  
12 }  
13 let foo_instance: foo = {  
14     field1 = {  
15         a = 10;  
16         b = { x= 2 ;y= 5; };  
17     };  
18     field2= 5;  
19 }
```

What's a root?

Think of them as root nodes of a graph, something that you have direct access to.

A **GC** will treat all the stack variables, registers, static section, as roots.

```
1  type baz = {  
2    x: int;  
3    y: int  
4  }  
5  type bar = {  
6    a: int;  
7    b: baz  
8  }  
9  type foo = {  
10   field1 : bar;  
11   field2: int  
12 }  
13 let foo_instance: foo = {  
14   field1 = {  
15     a = 10;  
16     b = { x= 2 ;y= 5; };  
17   };  
18   field2= 5;  
19 }
```

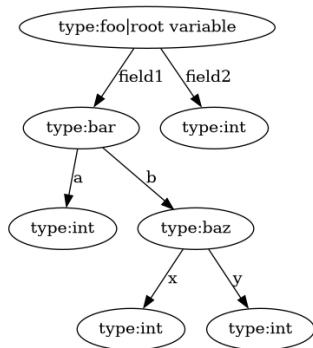


Figure 12: foo_instance variable as a graph

- It involves many intricate details, which took me quite a while to understand.

- It involves many intricate details, which took me quite a while to understand.
- You'd probably want to design a **precise-collector** if you're building your own language.

- It involves many intricate details, which took me quite a while to understand.
- You'd probably want to design a **precise-collector** if you're building your own language. **Precise Collector** is when GC knows about root objects: where they're placed on stack, their size, lifetime etc.

- It involves many intricate details, which took me quite a while to understand.
- You'd probably want to design a **precise-collector** if you're building your own language. **Precise Collector** is when GC knows about root objects: where they're placed on stack, their size, lifetime etc. Most popular language runtimes use precise garbage collectors: Python, PyPy, JVM(Java), .NET(C#), Lua, V8(JavaScript), SpiderMonkey(JavaScript) and a lot of others.

- It involves many intricate details, which took me quite a while to understand.
- You'd probably want to design a **precise-collector** if you're building your own language. **Precise Collector** is when GC knows about root objects: where they're placed on stack, their size, lifetime etc. Most popular language runtimes use precise garbage collectors: Python, PyPy, JVM(Java), .NET(C#), Lua, V8(JavaScript), SpiderMonkey(JavaScript) and a lot of others.
- Conservative collection does not know about types of traced GC objects.

- It involves many intricate details, which took me quite a while to understand.
- You'd probably want to design a **precise-collector** if you're building your own language. **Precise Collector** is when GC knows about root objects: where they're placed on stack, their size, lifetime etc. Most popular language runtimes use precise garbage collectors: Python, PyPy, JVM(Java), .NET(C#), Lua, V8(JavaScript), SpiderMonkey(JavaScript) and a lot of others.
- Conservative collection does not know about types of traced GC objects. It'll assume everything that's present in the stack, static and register are pointers and it'll follow it.

- It involves many intricate details, which took me quite a while to understand.
- You'd probably want to design a **precise-collector** if you're building your own language. **Precise Collector** is when GC knows about root objects: where they're placed on stack, their size, lifetime etc. Most popular language runtimes use precise garbage collectors: Python, PyPy, JVM(Java), .NET(C#), Lua, V8(JavaScript), SpiderMonkey(JavaScript) and a lot of others.
- Conservative collection does not know about types of traced GC objects. It'll assume everything that's present in the stack, static and register are pointers and it'll follow it. They are inherently unsafe.

- It involves many intricate details, which took me quite a while to understand.
- You'd probably want to design a **precise-collector** if you're building your own language. **Precise Collector** is when GC knows about root objects: where they're placed on stack, their size, lifetime etc. Most popular language runtimes use precise garbage collectors: Python, PyPy, JVM(Java), .NET(C#), Lua, V8(JavaScript), SpiderMonkey(JavaScript) and a lot of others.
- Conservative collection does not know about types of traced GC objects. It'll assume everything that's present in the stack, static and register are pointers and it'll follow it. They are inherently unsafe. Example: Boehm-Demers-Weiser conservative C/C++ Garbage Collector

All trace-based algorithms compute the set of reachable objects and then take the complement of this set. Memory is therefore recycled as follows:

All trace-based algorithms compute the set of reachable objects and then take the complement of this set. Memory is therefore recycled as follows:

1. The program or mutator runs and makes allocation requests.

All trace-based algorithms compute the set of reachable objects and then take the complement of this set. Memory is therefore recycled as follows:

1. The program or mutator runs and makes allocation requests.
2. The garbage collector discovers reachability by tracing.

All trace-based algorithms compute the set of reachable objects and then take the complement of this set. Memory is therefore recycled as follows:

1. The program or mutator runs and makes allocation requests.
2. The garbage collector discovers reachability by tracing.
3. The garbage collector reclaims the storage for unreachable objects.

All trace-based algorithms compute the set of reachable objects and then take the complement of this set. Memory is therefore recycled as follows:

1. The program or mutator runs and makes allocation requests.
2. The garbage collector discovers reachability by tracing.
3. The garbage collector reclaims the storage for unreachable objects.

Tracing Based GC Algorithms

All trace-based algorithms compute the set of reachable objects and then take the complement of this set. Memory is therefore recycled as follows:

1. The program or mutator runs and makes allocation requests.
2. The garbage collector discovers reachability by tracing.
3. The garbage collector reclaims the storage for unreachable objects.

Tracing Based GC Algorithms

- Copying Collector
- Mark and Sweep Collector

- Copying Collector partitions memory into two semispaces, A and B.

- Copying Collector partitions memory into two semispaces, A and B.
- All the allocations are performed in one semispace until it fills up.

- Copying Collector partitions memory into two semispaces, A and B.
- All the allocations are performed in one semispace until it fills up.
- The garbage collector then copies reachable objects to the other space.

- **Copying Collector** partitions memory into two semispaces, A and B.
- All the allocations are performed in one semispace until it fills up.
- The garbage collector then copies reachable objects to the other space.
- Roles of the semispaces are reversed when garbage collection is complete.

- **Copying Collector** partitions memory into two semispaces, A and B.
- All the allocations are performed in one semispace until it fills up.
- The garbage collector then copies reachable objects to the other space.
- Roles of the semispaces are reversed when garbage collection is complete.
- While all this is happening, the whole program is stopped. (This is why it's called a stop-the-world algorithm)

Copying Collector (A visual representation)




Figure 13: Before GC

Copying Collector (A visual representation)




Figure 13: Before GC




Figure 14: After GC

Some terminology first

Some terminology first

In a **mark-and-sweep** gc, An object may be in one of the 4 states: Free(Blue), Unreached(White), Unscanned(Gray), Scanned(Black)

Some terminology first

In a **mark-and-sweep** gc, An object may be in one of the 4 states: Free(Blue), Unreached(White), Unscanned(Gray), Scanned(Black)

It is also a stop-the-world algorithm.

Some terminology first

In a **mark-and-sweep** gc, An object may be in one of the 4 states: Free(Blue), Unreached(White), Unscanned(Gray), Scanned(Black)

It is also a stop-the-world algorithm.

Mark → responsible for marking reachable objects **Black**

Some terminology first

In a **mark-and-sweep** gc, An object may be in one of the 4 states: Free(Blue), Unreached(White), Unscanned(Gray), Scanned(Black)

It is also a stop-the-world algorithm.

Mark → responsible for marking reachable objects **Black**

Sweep → responsible for freeing all the unreachable(**White**) objects.

Mark and Sweep(continued)

```
/* marking phase */
1) add each object referenced by the root set to list Unscanned
   and set its reached-bit to 1;
2) while (Unscanned  $\neq \emptyset$ ) {
3)   remove some object o from Unscanned;
4)   for (each object o' referenced in o) {
5)     if (o' is unreachable; i.e., its reached-bit is 0) {
6)       set the reached-bit of o' to 1;
7)       put o' in Unscanned;
           }
       }
   }
/* sweeping phase */
8) Free =  $\emptyset$ ;
9) for (each chunk of memory o in the heap) {
10)  if (o is unreachable, i.e., its reached-bit is 0) add o to Free;
11)  else set the reached-bit of o to 0;
    }
```

Figure 15: Basic Mark and Sweep

Which is better ? (Automatic
Memory Management or Manual
Management)

Which is better ? (Automatic Memory Management or Manual Management)

Which is better ? (Automatic Memory Management or Manual Management)

DEPENDS

Case against Manual Memory Management

```
λ ./abc
[1] 55077 segmentation fault (core dumped) ./abc
temp on ↳ master [?] is 📦 v0.1.0 via 🐍 v3.10.10 via .
λ █
```

Figure 16: Segmentation Faults

Case against Manual Memory Management(continued)

Memory safety

The Chromium project finds that around 70% of our serious security bugs are [memory safety problems](#). Our next major project is to prevent such bugs at source.

The problem

Around 70% of our high severity security bugs are memory unsafety problems (that is, mistakes with C/C++ pointers). Half of those are use-after-free bugs.


High+, impacting stable

Security-related assert
7.1%

Other
23.9%

Other memory unsafety
32.9%

Use-after-free
36.1%



(Analysis based on 912 high or critical [severity](#) security bugs since 2015, affecting the Stable channel.)

These bugs are spread evenly across our codebase, and a high proportion of our non-security stability bugs share the same types of root cause. As well as risking our users' security, these bugs have real costs in how we fix and ship Chrome.

Figure 17: Use After Free

Case against Automatic Memory Management

We frequently performed an operation we called the “gossip dance”, where we’d take a node out of rotation to let it compact without taking traffic, bring it back in to pick up hints from Cassandra’s hinted handoff, and then repeat until the compaction backlog was empty. We also spent a large amount of time tuning the JVM’s garbage collector and heap settings, because GC pauses would cause significant latency spikes.

Figure 18: Discord Troubles(from a recent Discord Engineering Blog)

Case against Automatic Memory Management

We frequently performed an operation we called the “gossip dance”, where we’d take a node out of rotation to let it compact without taking traffic, bring it back in to pick up hints from Cassandra’s hinted handoff, and then repeat until the compaction backlog was empty. We also spent a large amount of time tuning the JVM’s garbage collector and heap settings, because GC pauses would cause significant latency spikes.

Figure 18: Discord Troubles(from a recent Discord Engineering Blog)

- Similar GC issues encountered by many tech giants

Advanced Topics in Memory Management

- Advanced type system like in Rust, which help with memory safety

- Advanced type system like in Rust, which help with memory safety
- Incremental GC

- Advanced type system like in Rust, which help with memory safety
- Incremental GC
- Parallel and Concurrent GC

- Advanced type system like in Rust, which help with memory safety
- Incremental GC
- Parallel and Concurrent GC
- Generational GC

- Advanced type system like in Rust, which help with memory safety
- Incremental GC
- Parallel and Concurrent GC
- Generational GC
- Reducing GC pauses

References

- <https://rebelky.cs.grinnell.edu/Courses/CS302/99S/Presentations/GC/>
- <https://discord.com/blog/how-discord-stores-trillions-of-messages>
- <https://gchandbook.org/>
- Compilers: Principles, Techniques, and Tools by Alfred V. Aho, Monica S. Lam, Ravi Sethi, and Jeffrey D. Ullman
- [cppreference](#)